

Handling Advanced Traffic Routing Patterns



Nigel Brown

@n_brownuk www.windsock.io



Module Outline



See how and why we might deploy multiple ingress controllers

Learn how to avoid contention between ingress controllers

Apply annotations to influence ingress controller behavior

Look at some advanced traffic routing patterns using ingress



Ingress API Summary




Ingress API

Provides the basics for managing ingress



Proxy Capabilities

Can we exploit reverse proxy features?



Annotations

Mechanism for applying metadata

Multiple Ingress Controllers



A cloud provider may already provide a default ingress controller



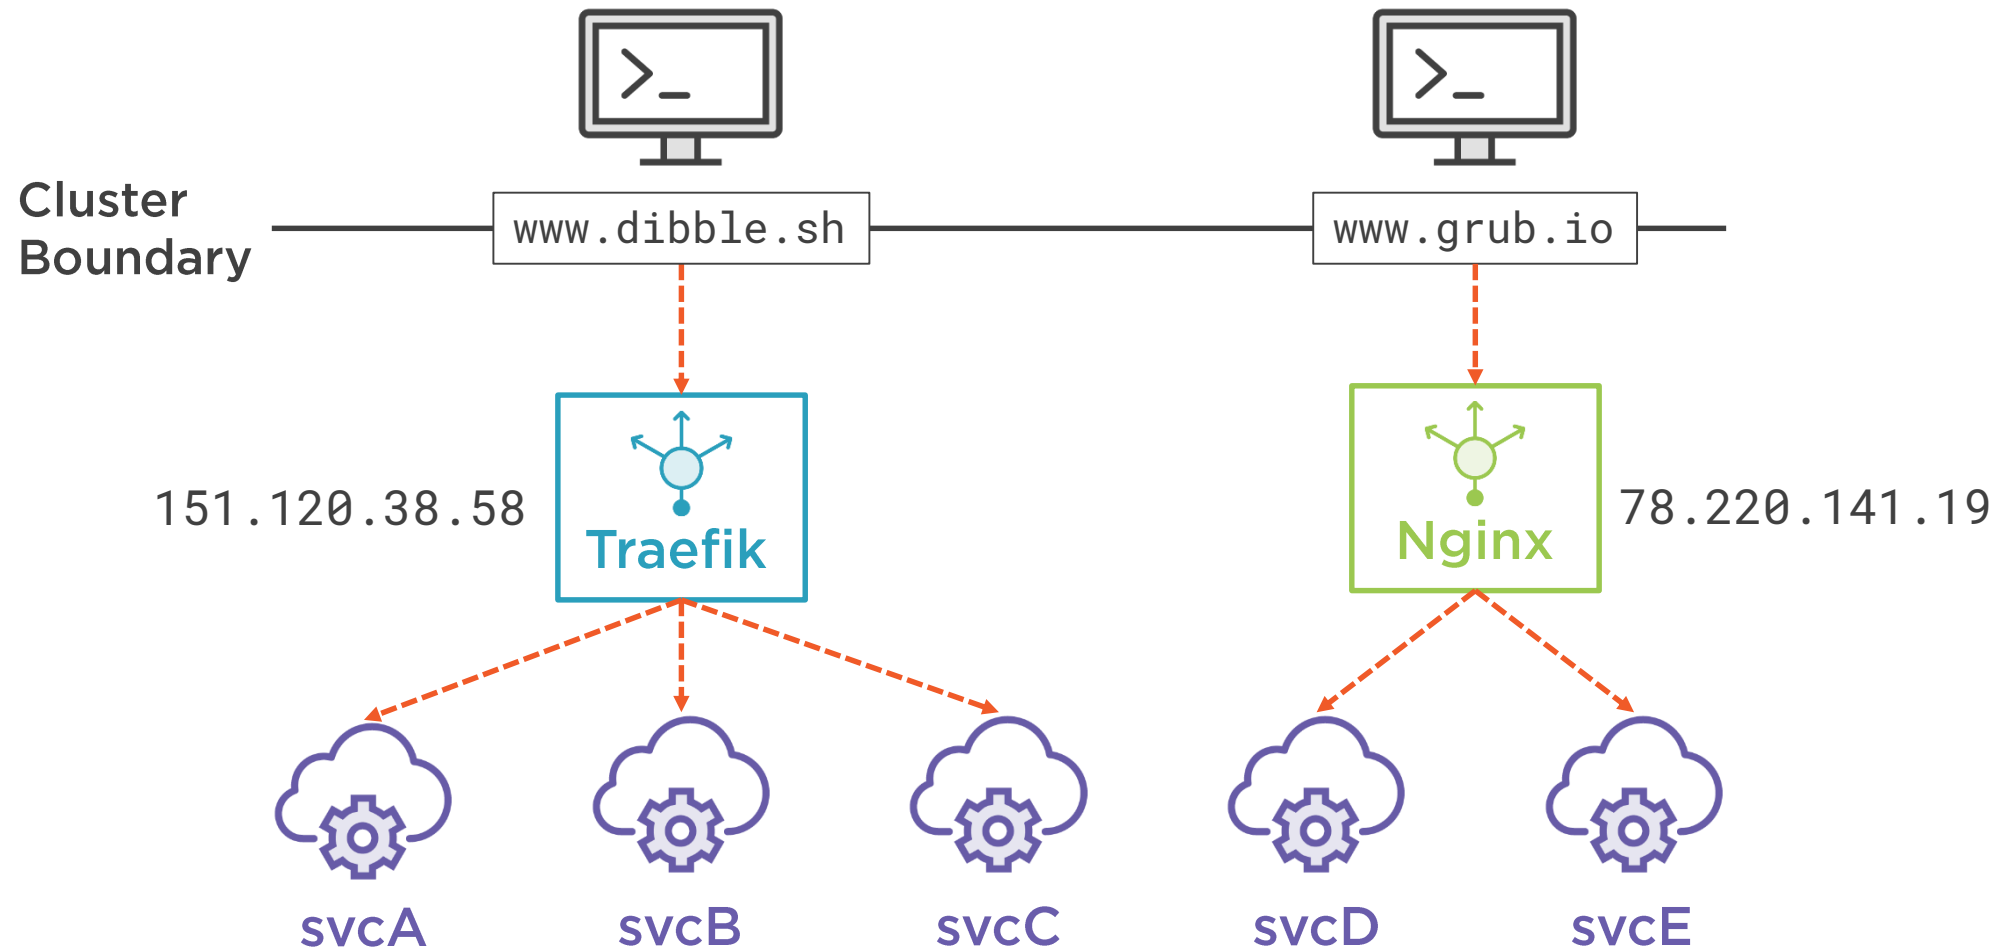
Perhaps traffic needs to be partitioned for security reasons



Not all ingress controllers provide the same functional capabilities



Deploying Multiple Ingress Controllers






Ingress Controller Contention

Competing ingress controllers will produce unpredictable traffic routing outcomes




Ingress Class



Nginx

Default class: nginx


`--ingress-class`



Traefik

Default class: traefik

`--kubernetes.ingressclass`



Contour

Default class: contour

`--ingress-class-name`



```
apiVersion: extensions/v1beta1
kind: Ingress
metadata
  name: contour-ingress
  annotations:
    kubernetes.io/ingress.class: "contour"
```

Assigning Ingress Objects to Controllers

A missing ingress class annotation may result in controller contention
It's good practice to annotate ingress definitions with an ingress class



Demo



Deploy two instances of the nginx ingress controller

One handles internal traffic, one handles external traffic

Individual backend services are deployed to serve both traffic types

Each ingress is configured with a class annotation



Canary Release

A deployment pattern, where a new release of software is made available to a small subset of users in order to minimize risk.



Canary Deployment




Both deployments share a **service** using a common **label**




Deployment Pattern Limitations




20% = 5 replicas



10% = 10 replicas



5% = 20 replicas



2% = 50 replicas



Reverse proxy load
balancing techniques can
weight traffic to upstream
virtual hosts



```
apiVersion: extensions/v1beta1
kind: Ingress
metadata
  name: traefik-ingress
  annotations:
    traefik.ingress.kubernetes.io/service-weights: |
      app: 90%
      app-canary: 10%
```

Weighting Traffic - Traefik Ingress Controller

The backend services must share the same host and path



```
apiVersion: extensions/v1beta1
kind: Ingress
metadata
  name: nginx-ingress
  annotations:
    nginx.ingress.kubernetes.io/canary: "true"
    nginx.ingress.kubernetes.io/canary-weight: "10"
```

Weighting Traffic - Nginx Ingress Controller

An association is made between the two different ingress definitions



Demo



Deploy an application service fronted by an ingress

Introduce a canary release of the application

Configure the ingress to send 5% of traffic to the canary

Test the ingress honors the specified traffic weighting



Module Summary



Limitations of Ingress API for advanced use cases

Annotations provide an effective workaround

Traffic segregation using ingress class annotations

Canary deployment pattern with ingress controller specific annotations

